

Importing Libraries

```
import pandas as pd
import numpy as np
```

Choose Dataset file from Local Directory

```
from google.colab import files
uploaded = files.upload()
```

Choose Files

Wine\_Quality\_Data.csv

- Wine\_Quality\_Data.csv(text/csv) - 457699 bytes, last modified: 3/26/2023 - 100% done

Saving Wine\_Quality\_Data.csv to Wine\_Quality\_Data.csv

Load Dataset

```
dataset = pd.read_csv('Wine_Quality_Data.csv')
```

```
print(dataset)
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	
...	...	...	...	...	...	
6492	6.2	0.21	0.29	1.6	0.039	
6493	6.6	0.32	0.36	8.0	0.047	
6494	6.5	0.24	0.19	1.2	0.041	
6495	5.5	0.29	0.30	1.1	0.022	
6496	6.0	0.21	0.38	0.8	0.020	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.99780	3.51	0.56	
1	25.0	67.0	0.99680	3.20	0.68	
2	15.0	54.0	0.99700	3.26	0.65	
3	17.0	60.0	0.99800	3.16	0.58	
4	11.0	34.0	0.99780	3.51	0.56	
...	...	...	...	...	...	
6492	24.0	92.0	0.99114	3.27	0.50	
6493	57.0	168.0	0.99490	3.15	0.46	
6494	30.0	111.0	0.99254	2.99	0.46	
6495	20.0	110.0	0.98869	3.34	0.38	
6496	22.0	98.0	0.98941	3.26	0.32	

	alcohol	quality	color
0	9.4	5	red
1	9.8	5	red
2	9.8	5	red
3	9.8	6	red
4	9.4	5	red
...	...	...	...
6492	11.2	6	white
6493	9.6	5	white
6494	9.4	6	white
6495	12.8	7	white
6496	11.8	6	white

[6497 rows x 13 columns]

Summarize Dataset

```
print(dataset.shape)
print(dataset.head(5))
```

(6497, 13)

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality	color
0	9.4	5	red
1	9.8	5	red
2	9.8	5	red
3	9.8	6	red
4	9.4	5	red

```

0      9.4      5  red
1      9.8      5  red
2      9.8      5  red
3      9.8      6  red
4      9.4      5  red

```

Covertng text to binary value

```

income_set = set(dataset['color'])
dataset['color'] = dataset['color'].map({'red':0, 'white':1}).astype(int)
print(dataset.head)

```

```

<bound method NDFrame.head of
0      7.4      0.70      0.00      1.9      0.076
1      7.8      0.88      0.00      2.6      0.098
2      7.8      0.76      0.04      2.3      0.092
3     11.2      0.28      0.56      1.9      0.075
4      7.4      0.70      0.00      1.9      0.076
...
6492      6.2      0.21      0.29      1.6      0.039
6493      6.6      0.32      0.36      8.0      0.047
6494      6.5      0.24      0.19      1.2      0.041
6495      5.5      0.29      0.30      1.1      0.022
6496      6.0      0.21      0.38      0.8      0.020

      free_sulfur_dioxide  total_sulfur_dioxide  density  pH  sulphates  \
0      11.0      34.0  0.99780  3.51      0.56
1      25.0      67.0  0.99680  3.20      0.68
2      15.0      54.0  0.99700  3.26      0.65
3      17.0      60.0  0.99800  3.16      0.58
4      11.0      34.0  0.99780  3.51      0.56
...
6492      24.0      92.0  0.99114  3.27      0.50
6493      57.0     168.0  0.99490  3.15      0.46
6494      30.0     111.0  0.99254  2.99      0.46
6495      20.0     110.0  0.98869  3.34      0.38
6496      22.0      98.0  0.98941  3.26      0.32

      alcohol  quality  color
0      9.4      5      0
1      9.8      5      0
2      9.8      5      0
3      9.8      6      0
4      9.4      5      0
...
6492     11.2      6      1
6493      9.6      5      1
6494      9.4      6      1
6495     12.8      7      1
6496     11.8      6      1

[6497 rows x 13 columns]>

```

Segragate Dataset int X & Y

```

X = dataset.iloc[:, :-1].values
X

```

```

array([[ 7.4 ,  0.7 ,  0. , ...,  0.56,  9.4 ,  5. ],
       [ 7.8 ,  0.88,  0. , ...,  0.68,  9.8 ,  5. ],
       [ 7.8 ,  0.76,  0.04, ...,  0.65,  9.8 ,  5. ],
       ...,
       [ 6.5 ,  0.24,  0.19, ...,  0.46,  9.4 ,  6. ],
       [ 5.5 ,  0.29,  0.3 , ...,  0.38, 12.8 ,  7. ],
       [ 6. ,  0.21,  0.38, ...,  0.32, 11.8 ,  6. ]])

```

```

Y = dataset.iloc[:, -1].values
Y

```

```

array([0, 0, 0, ..., 1, 1, 1])

```

Splitting Dataset into Train & Test

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.25, random_state = 0)

```

Feature Scaling

```

from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)

```

```
X_test = sc.transform(X_test)
print(X_train)
```

```
[[[-0.08802117 -0.72820076  0.00517689 ... -1.00933084 -0.75281497
   0.20911623]
 [-0.78538358 -1.09072305  0.07390119 ... -0.47886052  0.34073703
   1.36031877]
 [-0.39796002 -0.60736   -0.13227171 ...  0.38315375  0.67721458
   2.51152131]
 ...
 [-0.32047531 -0.84904152  1.86073296 ... -1.00933084 -1.1734119
   0.20911623]
 [-0.47544473 -0.72820076  0.34879839 ...  0.11791859 -0.16397928
   1.36031877]
 [ 1.15173422 -0.72820076  0.69241988 ... -0.67778689 -0.41633743
  -0.94208632]]
```

## Training

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression() # loading the algorithm
model.fit(X_train, y_train) # train
```

```
▼ LogisticRegression
LogisticRegression()
```

## Predict for all Test Data

```
y_pred = model.predict(X_test)
```

## Evaluating the model

```
from sklearn.metrics import accuracy_score
print("Accuracy of the Model: {}".format(accuracy_score(y_test, y_pred)*100))
```

```
Accuracy of the Model: 99.44615384615385%
```

## Confusion Matrix

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix: ")
print(cm)
```

```
Confusion Matrix:
[[ 388   4]
 [   5 1228]]
```

## Predicting whether the wine is red or white

```
# fixed_acidity = float(input("Enter fixed_acidity: "))
# volatile_acidity = float(input("Enter volatile_acidity: "))
# citric_acid = float(input("Enter citric_acid: "))
# residual_sugar = float(input("Enter residual_sugar: "))
# chlorides = float(input("Enter chlorides: "))
# free_sulfur_dioxide = float(input("Enter free_sulfur_dioxide: "))
# total_sulfur_dioxide = float(input("Enter total_sulfur_dioxide: "))
# density = float(input("Enter density: "))
# pH = float(input("Enter pH: "))
# sulphates = float(input("Enter sulphates: "))
# alcohol = float(input("Enter alcohol: "))
# quality = int(input("Enter quality: "))
# newWine = [[fixed_acidity, volatile_acidity, citric_acid, residual_sugar,
#             chlorides, free_sulfur_dioxide, total_sulfur_dioxide, density,
#             pH, sulphates, alcohol, quality]]

# result = model.predict(sc.transform(newWine))

# print(result)

# if result == 'red':
#     print("The wine is red.")
# else:
#     print("The wine is white")
```

```
Enter fixed_acidity: 6.5
Enter volatile_acidity: 0.22
```

```

Enter citric_acid: 0.5
Enter residual_sugar: 16.4
Enter chlorides: 0.048
Enter free_sulfur_dioxide: 36
Enter total_sulfur_dioxide: 182
Enter density: 0.99904
Enter pH: 3.02
Enter sulphates: 0.49
Enter alcohol: 8.8
Enter quality: 6
['white']
The wine is white

```

### Receiver Operating Curve - ROC Curve


```

from sklearn.metrics import roc_auc_score, roc_curve
import matplotlib.pyplot as plt

nsProbability = [0 for _ in range(len(y_test))]
lsProbability = model.predict_proba(X_test)
# keep probabilities for the positive outcome only
lsProbability = lsProbability[:, 1]
# calculate score
nsAUC = roc_auc_score(y_test, nsProbability)
lsAUC = roc_auc_score(y_test, lsProbability)
# summarize score
print("NO skill: ROC AUC=%.3f" % (nsAUC*100))
print("Logistic skill: ROC AUC=%.3f" % (lsAUC*100))
# calculate roc curves
nsFP, nsTP, _ = roc_curve(y_test, nsProbability)
lrFP, lrTP, _ = roc_curve(y_test, lsProbability)
# plot the roc curve for the model
plt.plot(nsFP, nsTP, linestyle='--', label='No Skill')
plt.plot(lrFP, lrTP, marker='*', label='Logistic')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
# show the legend
plt.legend()
plt.show()

```

NO skill: ROC AUC=50.000  
Logistic skill: ROC AUC=99.907



### Cross Validation Score

```

from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
kfold = KFold(n_splits=10)
result = cross_val_score(model, X, Y, cv=kfold)
print("CROSS VALIDATION SCORE: %.2f%%" % (result.mean()*100.0))

```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
CROSS VALIDATION SCORE: 97.23%
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

## Stratified K-fold Cross Validation

```
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import StratifiedKFold
skfold = StratifiedKFold(n_splits=3)
model_skfold = LogisticRegression()
results_skfold = cross_val_score(model_skfold, X, Y, cv=skfold)
print("STRATIFIED K-FOLD SCORE: %.2f%%" % (results_skfold.mean()*100.0))
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
STRATIFIED K-FOLD SCORE: 97.86%
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

## Cumulative Accuracy Profile (CAP) Curve

```
total = len(y_test)
class_1_count = np.sum(y_test)
```

```

print(class_1_count)
class_0_count = total - class_1_count
plt.plot([0, total], [0, class_1_count], c='r', linestyle='--', label='Random Model')
plt.plot([0, class_1_count, total],
         [0, class_1_count, class_1_count],
         c='grey',
         linewidth=2,
         label='Perfect Model')
probs = model.predict_proba(X_test)
probs = probs[:, 1]
model_y = [y for _, y in sorted(zip(probs, y_test), reverse=True)]
y_values = np.append([0], np.cumsum(model_y))
x_values = np.arange(0, total+1)

plt.plot(x_values, y_values, c='b', label='LR Classifier', linewidth=4)
index = int((50*total / 100))

## 50% Vertical line from x-axis
plt.plot([index, index], [0, y_values[index]], c='g', linestyle='--')

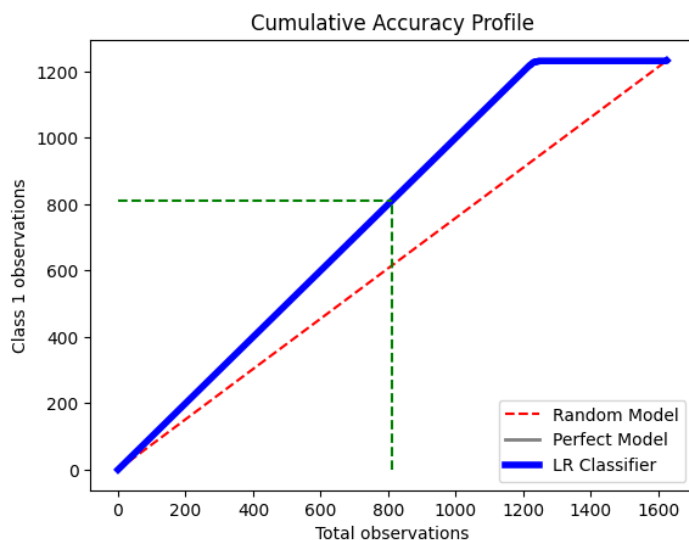
## Horizontal line to y-axis from prediction model
plt.plot([0, index], [y_values[index], y_values[index]], c='g', linestyle='--')

class_1_observed = y_values[index] * 100 / max(y_values)
plt.xlabel('Total observations')
plt.ylabel('Class 1 observations')
plt.title('Cumulative Accuracy Profile')
plt.legend(loc = 'lower right')

```

1233

&lt;matplotlib.legend.Legend at 0x7f39b330a940&gt;



✓ 0s completed at 6:18 AM

● ×